

Project 5

Benjamin Draves

```
#-----
# Neural Network Metropolis-within-Gibbs Sampler
#-----

# helper functions
LOGEPS <- log(.Machine$double.eps/2)
log1pe <- function(x) {
  # vectorized version: `x` can be a vector
  l <- ifelse(x > 0, x, 0) # shift
  x <- ifelse(x > 0, -x, x) # range reduction: `x = -abs(x)`
  ifelse(x < LOGEPS, 1, 1 + log(1 + exp(x)))
}

# ft: the neural network 'function' - wrapper function
# arguments: 1. Data vector x 2. Model parameter vector a0 2.
# Model parameter matrix a (ak in the rows of a) 2. Model
# parameter vector b
ft <- function(x, a0, a, b) {
  K <- nrow(a)
  tot <- 0
  for (k in 1:K) {
    tot <- tot + b[k] * tanh(a0[k] + sum(a[, k] * x))
  }
  return(tot)
}

# f: the full neural network - wrapper function arguments: 1.
# Data matrix x 2. Model parameter vector a0 2. Model parameter
# matrix a (ak in the rows of a) 2. Model parameter vector b
f <- function(X, a0, a, b) apply(X, 1, function(x) ft(x, a0,
  a, b))

# logpi: A function that evaluates log[pi(*|theta\j, Y, X)]
# arguments: 1. Response vector Y 2. Design matrix X 3. Fixed
# standard deviation c^2 4. The current parameter value theta
# 5. Parameter values (a0, a, b)

logpi <- function(Y, X, c2, theta, a0, a, b) {
  eval <- f(X, a0, a, b)

  if (sum(is.na(eval)) != 0) {
    print("STOPPED AT logpi")
    break
  }

  -(theta)^2/(2 * c2) + sum(Y * eval) - sum(log1pe(eval))
}

# gibbs_sampler: a function that returns the samples from the
# MCMC arguments: 1. Response vector Y 2. Design matrix X 3.
# Initial parameter value matrix theta0 (K, 2 + p) - (a_0k,
# a_k, b_k) 4. K - depth of NN (default K = 10) 5. Initial
# 'step sizes' matrix v (K, p + 2) 6. Learning rate function
# gamma: t --> 1/t^gamma (default = 2) 7. Desired acceptance
# rate r (default = 30%) 8. Maximum number of iterations
# (default = 5000)
```

```

gibbs_sampler <- function(Y, X, theta0, v, K = 10, gamma = 2,
  c2 = 1, r = 0.3, max.iter = 5000) {

  #-----
  # set up data structures - paramter values
  #-----
  p <- ncol(X)
  tfm <- NA

  # Intercept term
  a0 <- matrix(NA, nrow = K, ncol = max.iter)
  a0[, 1] <- theta0[, 1]

  # Weights
  b <- matrix(NA, nrow = K, ncol = max.iter)
  b[, 1] <- theta0[, p + 2]

  # Slope terms
  a <- array(NA, dim = c(K, p, max.iter))
  a[, , 1] <- theta0[, 2:(p + 1)]

  #-----
  # set up data structures - step size
  #-----
  v0 <- matrix(NA, nrow = K, ncol = max.iter)
  v0[, 1] <- v[, 1]

  vb <- matrix(NA, nrow = K, ncol = max.iter)
  vb[, 1] <- v[, p + 2]

  va <- array(NA, dim = c(K, p, max.iter + 1))
  va[, , 1] <- v[, 2:(p + 1)]

  #-----
  # set up data structures - acceptance 0/1
  #-----
  acceptance <- array(NA, dim = c(K, p + 2, max.iter))

  # begin samples
  for (t in 2:max.iter) {
    #-----
    # Update parameters
    #-----

    # copy over paramaters for efficiency
    a0[1:K, t] <- a0[1:K, t - 1]
    b[1:K, t] <- b[1:K, t - 1]
    a[, , t] <- a[, , t - 1]

    # update a0
    for (k in 1:K) {
      # propose new sample  $N(a0[k, t-1], v0[k, t-1])$ 
      tstar <- rnorm(1, a0[k, t - 1], v0[k, t - 1])

      # calculate log acceptance probability
      top <- a0[1:K, t]
      top[k] <- tstar
      logp <- min(c(0, logpi(Y, X, c2, tstar, top, a[,
        , t], b[, t]) - logpi(Y, X, c2, a0[k, t], a0[,
        t], a[, , t], b[, t])))
    }
  }
}

```

```

# Update new value
a0[k, t] <- ifelse(log(runif(1)) < logp, tstar, a0[k,
t])

# Update step size
v0[k, t] <- v0[k, t - 1] + 1/(t^(gamma)) * (exp(logp) -
r)

# update acceptance
acceptance[k, 1, t - 1] <- ifelse(tstar == a0[k,
t], 1, 0)
}

# update b
for (k in 1:K) {
# propose new sample  $N(b[k, i, t-1], vb[k, i, t-1])$ 
tstar <- rnorm(1, b[k, t - 1], vb[k, t - 1])

# calculate log acceptance probability
top <- b[1:K, t]
top[k] <- tstar
logp <- min(c(0, logpi(Y, X, c2, tstar, a0[, t],
a[, , t], top) - logpi(Y, X, c2, b[k, t], a0[,
t], a[, , t], b[, t])))

# Update new value
b[k, t] <- ifelse(log(runif(1)) < logp, tstar, b[k,
t])

# Update step size
vb[k, t] <- vb[k, t - 1] + 1/(t^(gamma)) * (exp(logp) -
r)

# update acceptance
acceptance[k, p + 2, t - 1] <- ifelse(tstar == b[k,
t], 1, 0)
}

# update a
for (k in 1:K) {
for (i in 1:p) {
# propose new sample  $N(a[k, i, t-1], va[k, i, t-1])$ 
tstar <- rnorm(1, a[k, i, t - 1], va[k, i, t -
1])

# calculate log acceptance probability
top <- a[, , t]
top[k, i] <- tstar
logp <- min(c(0, logpi(Y, X, c2, tstar, a0[,
t], top, b[, t]) - logpi(Y, X, c2, a[k, i,
t - 1], a0[, t], a[, , t], b[, t])))

# Update new value
a[k, i, t] <- ifelse(log(runif(1)) < logp, tstar,
a[k, i, t - 1])

# Update step size
va[k, i, t] <- va[k, i, t - 1] + 1/(t^(gamma)) *
(exp(logp) - r)

```

```

        # update acceptance
        acceptance[k, i + 1, t - 1] <- ifelse(tstar ==
        a[k, i, t], 1, 0)
    }
}

#-----
# Check Convergence
#-----
print(t)
if (t == max.iter) {
    tfin <- max.iter
    break
}

}

# format results
res <- array(NA, dim = c(K, p + 2, tfin))
res[, 1, ] <- a0
res[, 2:(p + 1), ] <- a
res[, p + 2, ] <- b

# return(list(parameters = res, acceptance = acceptance))
return(parameters = res)

}

#-----
# Parallel Tempering algorithm
#-----
# update_paramter: function that completes an update for PT
# parameter arguments: 1. Y: response vector (n x 1) 2. X:
# covariate matrix (n x p) 3. c2: numeric prior variance 4.
# theta: current value of parameters in each chain (d x 1)
# 5. params: most recent parameters of full model (K x p+2)
# 6. vec: which element we are updating in params (2 x 1) 7.
# temps: vector of temperatures (d x 1) 8. stepsize: numeric
# step size

update_parameter <- function(Y, X, c2, theta, params, vec, temps,
    step_size) {

    # set up local variables
    d <- length(temps)
    p <- ncol(X)
    theta_bar <- numeric(d)
    lprobs <- numeric(d)

    for (ell in 1:d) {
        # update each chain
        tstar <- rnorm(1, mean = theta[ell], sd = step_size)

        # calculate log acceptance probability
        top <- params
        top[vec[1], vec[2]] <- tstar
        logp <- 1/temps[ell] * min(c(0, logpi(Y, X, c2, tstar,
            top[, 1], top[, 2:(p + 1)], top[, p + 2]) - logpi(Y,
            X, c2, theta[ell], params[, 1], params[, 2:(p + 1)],

```

```

        params[, 2 + p]))

    # store log acceptance prob
    lprobs[ell] <- logp

    # update estimate
    theta_bar[ell] <- ifelse(log(runif(1)) < logp, tstar,
        theta[ell])
}

# draw I in (1, d-1)
I <- sample(1:(d - 1), 1)

# calculate paramter matrix with theta_bar[I]/theta_bar[I+1]
top <- params
top[vec[1], vec[2]] <- theta_bar[I]
bottom <- params
bottom[vec[1], vec[2]] <- theta_bar[I + 1]

# caculate log alpha
diff_h <- logpi(Y, X, c2, theta_bar[I], top[, 1], top[, 2:(p +
    1)], top[, p + 2]) - logpi(Y, X, c2, theta_bar[I + 1],
    bottom[, 1], bottom[, 2:(p + 1)], bottom[, p + 2])
logalpha <- min(c(0, (1/(temps[I + 1]) - 1/(temps[I])) *
    diff_h))

# swap /no swap
if (log(runif(1)) < logalpha) {
    # if yes swap
    tmp <- theta_bar[I]
    theta_bar[I] <- theta_bar[I + 1]
    theta_bar[I + 1] <- tmp
}

# return update values of each chain
return(list(parameter = theta_bar[1], positions = theta_bar[-1],
    lprobs = lprobs))
}

# update_paramter: function that completes an update for PT
# parameter arguments: 1. Y: response vector (n x 1) 2. X:
# covariate matrix (n x p) 3. theta0: initial parameter value
# matrix (K x p+2) 4. v0: initial step size matrix (K x p+2)
# 5. temps: initial temperatures vector (d x 1) 6. K: neural
# network depth K 7. c2: prior variance, default 1 8.
# max.iter: maximum number of iterations, defaul 5000 Used
# for adaptive step size 9. gamma: learning rate (.5, infty),
# default INFTY (no adaptive updates) 10. r: desired
# acceptance probability, default .3
pt_sampler <- function(Y, X, theta0, v0, temps, K = 10, c2 = 1,
    max.iter = 5000, gamma = Inf, r = 0.3) {

    # set up local variables
    p <- ncol(X)
    d <- length(temps)

    # set up storage for parallel chains
    pt_chains_positions <- array(1, dim = c(K, p + 2, d - 1))

```

```

# set up storage for parameter vectors
theta <- array(1, dim = c(K, p + 2, max.iter))
theta[, , 1] <- theta0

# set up storage for stepsizes
v <- array(1, dim = c(K, p + 2, max.iter))
v[, , 1] <- v0

for (t in 2:max.iter) {

  # copy over past parameters for computational convience
  theta[, , t] <- theta[, , t - 1]

  for (k in 1:K) {
    for (j in 1:(p + 2)) {
      # get current placement of chains
      theta_here <- c(theta[k, j, t], pt_chains_positions[k,
        j, ])

      # update parameters
      tmp <- update_parameter(Y, X, c2, theta_here,
        theta[, , t], c(k, j), temps, v[k, j, t - 1])
      theta[k, j, t] <- tmp$parameter
      pt_chains_positions[k, j, ] <- tmp$positions

      # update step size
      v[k, j, t] <- v[k, j, t - 1]  ## 1/(t^(gamma)) * (exp(tmp$lprobs[1]) - r)

    }
  }
  print(t)
}

return(theta)

```

```

}

```